

John Doe

Senior Data Engineer · 7 YOY · AWS Certified Data Analytics Specialist

+1 (555) 000-0000 · john.doe@example.com
johndoe.dev · github.com/johndoe · linkedin.com/in/johndoe

Data engineer with 7 years of experience delivering production-ready data platforms across life sciences, retail, SaaS, and enterprise intelligence domains. Specialized in end-to-end lakehouse architecture, real-time streaming pipelines, and agentic BI systems on AWS, with full CI/CD ownership from design through production.

EXPERIENCE

Tech Solutions Inc.

Nov 2022

Senior Data Engineer

- Delivered production-ready data platform programs for life sciences and retail clients, owning design, implementation, testing, CI/CD, and support across DEV, QA, and PROD environments.
- Architected an enterprise Customer Data Platform on AWS, unifying multiple data sources and legacy systems into a governed lakehouse and warehouse foundation serving 4 downstream applications.
- Designed an event-driven migration factory (Step Functions + Lambda + S3 events) that automated a 10 TB+ data warehouse migration, reducing end-to-end migration time by 70%.
- Led design and delivery of a two-plane NL2SQL data agent platform integrating LLM APIs with a governed semantic layer, vector search, and React / TypeScript front-end.

AWS · Python · Redshift · dbt · Databricks · Terraform · Step Functions · Lambda · FastAPI · CI/CD · Docker

SaaS Innovation Ltd.

Jun 2021 – Nov 2022

Data Engineer

- Designed and operated real-time and batch data platforms for enterprise SaaS products, serving millions of enterprises and 180 million global records.
- Built streaming ETL pipelines with Kafka and Flink; managed lakehouse models across multiple systems and AWS S3.
- Drove Agile delivery as a core team member across bi-weekly sprints, consistently meeting velocity targets.

Python · Spark · Flink · Kafka · TiDB · AWS · Hive · CI/CD · Docker · Agile

Data Systems Corp.

Jun 2019 – Jun 2021

Data Engineer

- Owned data engineering, building ETL pipelines and data models around 230 million+ enterprise entities.
- Improved data freshness from T+1 to near-real-time via optimized Python and Hive batch pipelines.
- Built and maintained search indexes supporting 100M+ document full-text search for enterprise risk signal detection.

Python · Hive · Airflow · MySQL · Elasticsearch · Hadoop · ETL

EDUCATION

State University of Technology

Sep 2015 – Jul 2019

Bachelor of Engineering · Computer Science and Technology

- Third Prize — National University Cloud Computing Application Innovation Competition.
- GPA 3.6 / 4.0, ranked top 15% of cohort.

PROJECTS

Enterprise Customer Data Platform

Dec 2022

End-to-end AWS customer data platform with governed ingestion, warehouse foundations, and secure data APIs.

- Built from 0 to 1 on AWS, unifying multiple data sources into a governed lakehouse; delivered column-level masking, row-level security, and automated quality gates via Great Expectations + dbt tests.
- Exposed governed data via FastAPI + API Gateway REST APIs consumed by 4 downstream client applications in production.

AWS · Redshift · dbt · Glue · S3 · FastAPI · Terraform · Great Expectations · Python

Enterprise Data Warehouse Migration

Nov 2022 – May 2023

Automated migration of 10 TB+ data warehouse to cloud-native architecture.

- Designed an event-driven migration factory with Step Functions, Lambda, S3 events, and staging databases — turning a fully manual process into a one-click repeatable pipeline with 70% time reduction and zero data loss.

AWS · Step Functions · Lambda · RDS · Redshift · Python · Event-Driven

Real-time Analytics Platform

Aug 2023 – Feb 2024

High-throughput streaming data pipeline processing 500M+ events daily with sub-second latency.

- Built end-to-end streaming architecture using Kafka, Flink, and Databricks Delta Lake, achieving 99.99% uptime and sub-100ms p99 latency for real-time dashboards.
- Implemented auto-scaling policies reducing infrastructure costs by 35% while maintaining performance SLAs.

Kafka · Apache Flink · Databricks · Delta Lake · Python · AWS · Prometheus · Grafana

Enterprise Data Agent Platform

Feb 2026

Full-stack agentic BI platform with governed semantics, NL2SQL, and interactive analytics.

- Led end-to-end architecture of a two-plane system decoupling semantic governance from runtime execution, enabling business users to query data via natural language with automatic validation and explainability.

Python · LLM APIs · FastAPI · TypeScript · React · PostgreSQL · Vector Search · AWS · NL2SQL

Data Governance & Compliance Framework

Mar 2024 – Sep 2024

Centralized metadata and governance system for enterprise data estate.

- Designed and implemented data contracts, automated PII detection, and audit logging serving 50+ data producers and 200+ consumers.
- Reduced compliance violations by 92% and data incident response time from 4 hours to 15 minutes through real-time monitoring.

Python · Apache Atlas · Data Contracts · dbt · PostgreSQL · Kubernetes · CI/CD

SKILLS

Data Platform Engineering	(Expert)	Python · SQL · dbt · Redshift · Snowflake · Databricks · Hive · Delta Lake · Apache Iceberg
Orchestration, Streaming & Compute	(Advanced)	Apache Airflow · Apache Spark · Apache Flink · Kafka · AWS Step Functions · AWS Glue · EMR
Cloud & Infrastructure	(Advanced)	AWS (S3, EC2, Lambda, ECS, RDS, Glue, Redshift) · Terraform · Docker · CI/CD · GitHub Actions · Kubernetes
Quality, Security & Reliability	(Advanced)	Great Expectations · dbt tests · Column-level Masking · Row-level Security · Data Contracts · Compliance
APIs & Application Layer	(Advanced)	FastAPI · REST / OpenAPI · TypeScript · React · LLM APIs · gRPC
Delivery & Collaboration	(Advanced)	Agile / Scrum · Technical Design · Stakeholder Management · Cross-functional Leadership · Mentoring

CERTIFICATIONS

AWS Certified Data Analytics – Specialty	Amazon Web Services · Jan 2024
AWS Certified Solutions Architect – Associate	Amazon Web Services · Mar 2023
Apache Spark Certified Developer	Databricks · Nov 2022
Kubernetes Application Developer	Cloud Native Computing Foundation · Jun 2023

LANGUAGES & INTERESTS

LANGUAGES	English (Native) · Mandarin Chinese (Professional Working) · Japanese (Elementary)
INTERESTS	Open Source (data tooling, Apache community, contributor) · Technical Writing (architecture blogs, Medium publications, conference talks) · Community (mentoring junior engineers, conference speaking, knowledge sharing)